

Project 1 - AWS SAA-C03

Scenario:

You have been asked to set up a WordPress blog for your company per their defined specifications.

During business hours (9AM-6PM), the blogging team would like to use their own development copy of the WordPress instance, so that any work they are doing doesn't impact the live copy.

You are required to create a plan to achieve these goals using Amazon CloudFormation, creating a new AMI and Auto Scaling.

Key elements of the project:

Use Amazon CloudFormation to create an EC2 instance to run WordPress with the following specification:

- Instance Type: T2.micro
- Create a new AMI of the WordPress instance.
- Configure Auto Scaling to launch a new WordPress instance during 9AM-6PM.

Solution:

To achieve the goal of setting up a WordPress blog with a development environment using Amazon CloudFormation, AMI creation, and Auto Scaling, here's a detailed plan:

First I create an instance for wordpress live copy:

1. Select "CloudFormation" from the dashboard and then click on "Create New Stack".
2. Then I select use a sample template and from the dropdown of sample templates and choose "WordPress blog" and click Next.
3. I enter a name (**Project-01**) in the "Stack Name" box and complete the rest of the options like username and DB password etc and change the Instance Type to t2 micro.
4. In the KeyName section, I select **mytest-key** that I created before in **Ohio region**. Then click next and Review the settings and click "Create". Wait until the Stack process has a status of "COMPLETED".

Below are the screenshots of my AWS console.

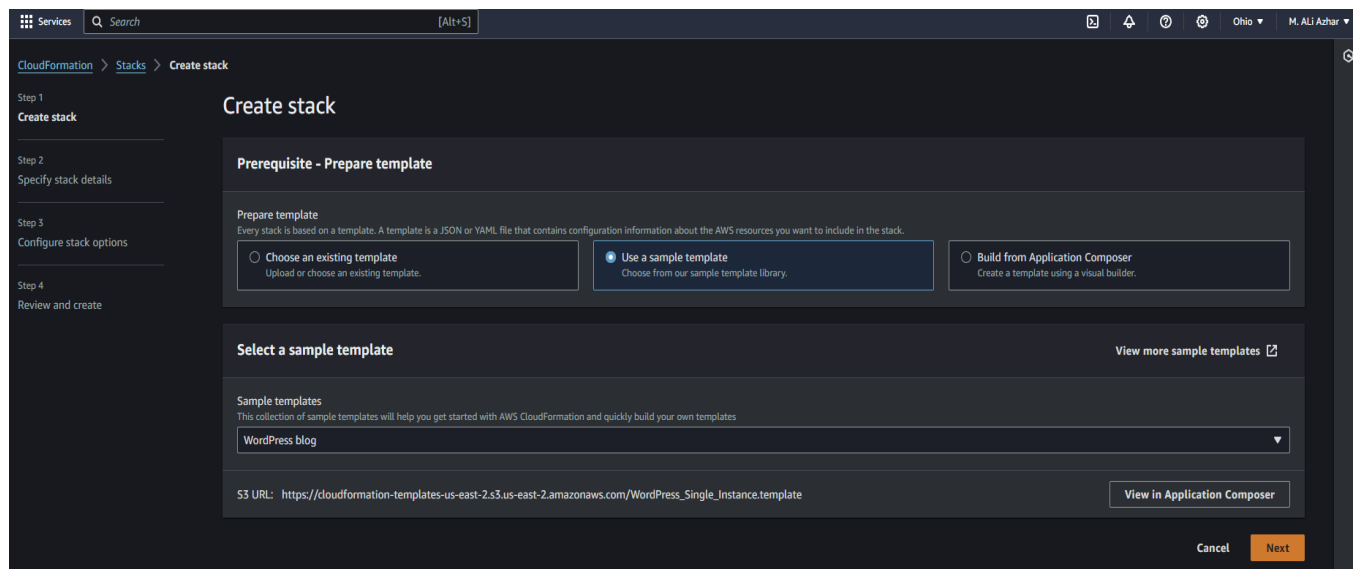


Fig 1: Cloud formation Stack Creation

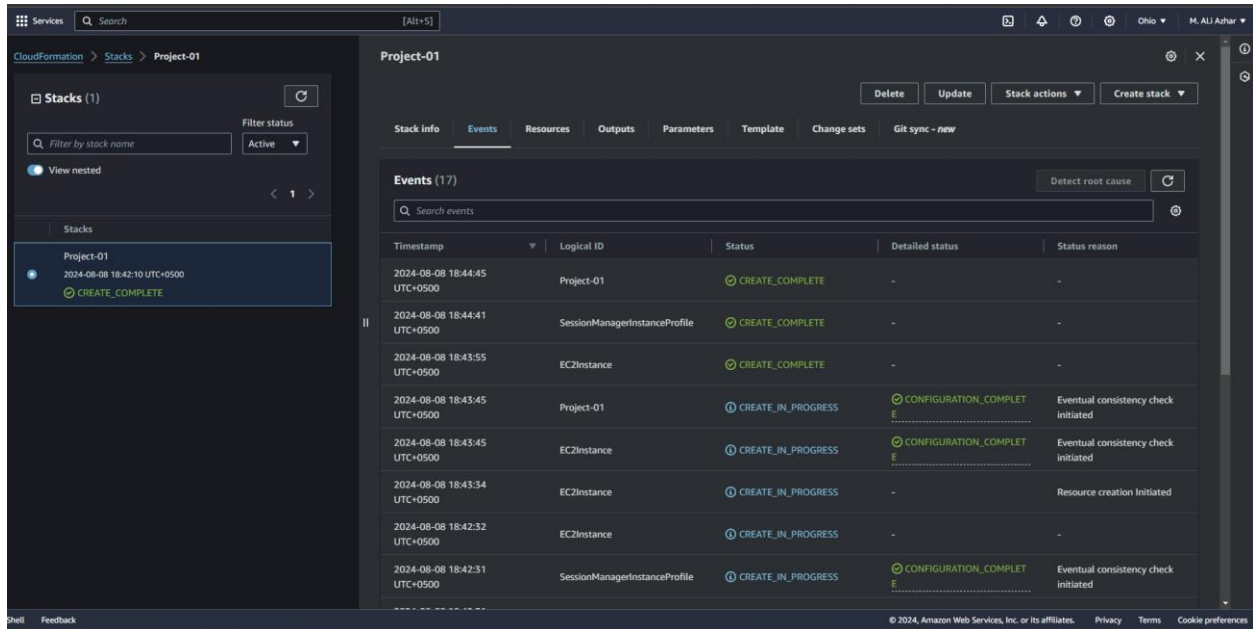


Fig 4: Project-01 stack status completed

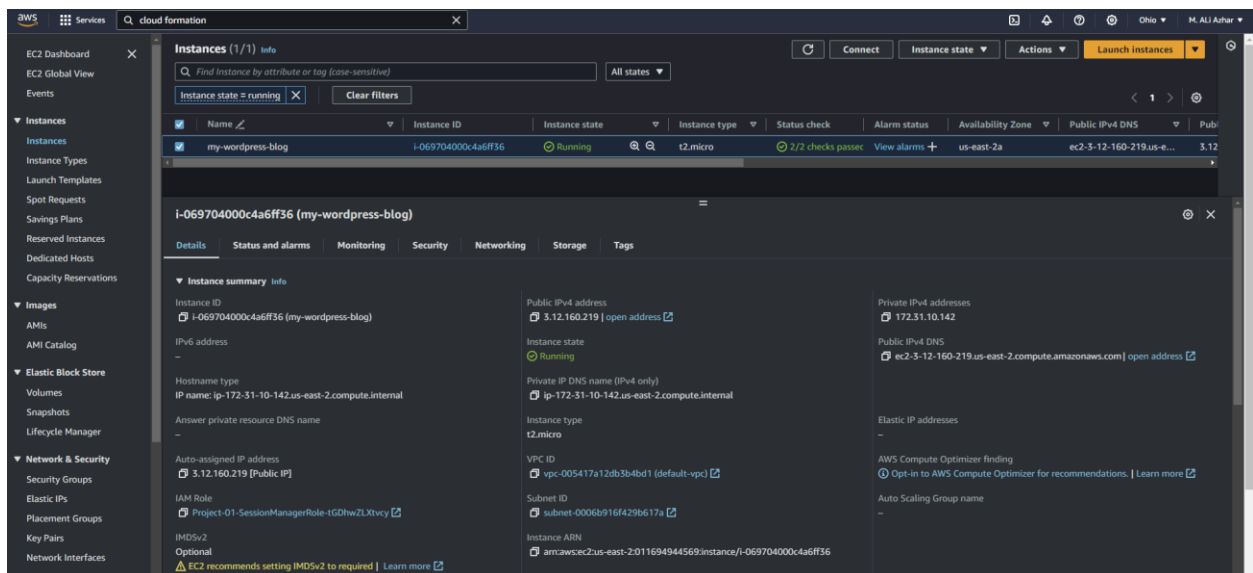


Fig 5: Wordpress Ec2 instance created from cloudformation Stack

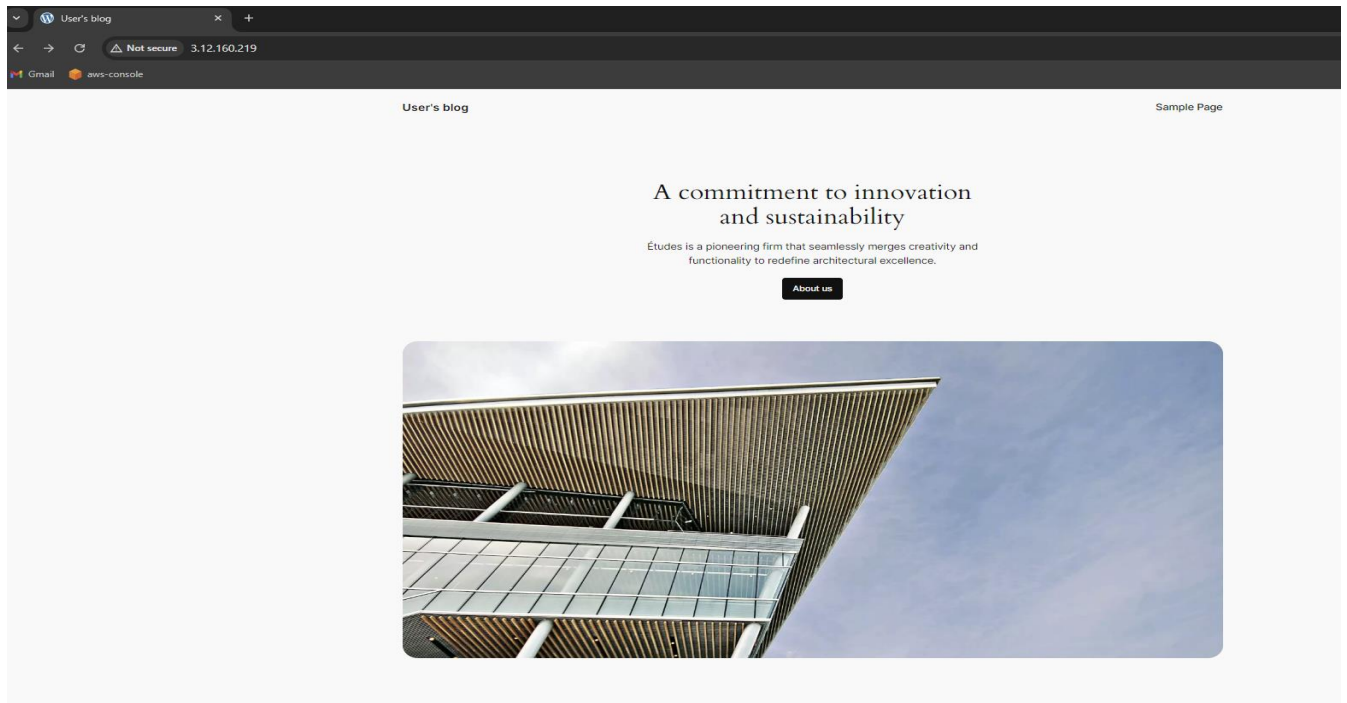


Fig 6: Wordpress sample user blog page

AMI Creation:

- I select the newly created instance and click on Actions > Image > Create Image.
- Here I enter a name (**my-wordpress-blog**) for my AMI image and then click on Create Image.
- I switch to the AMI dashboard and wait until my new AMI has a status of available.

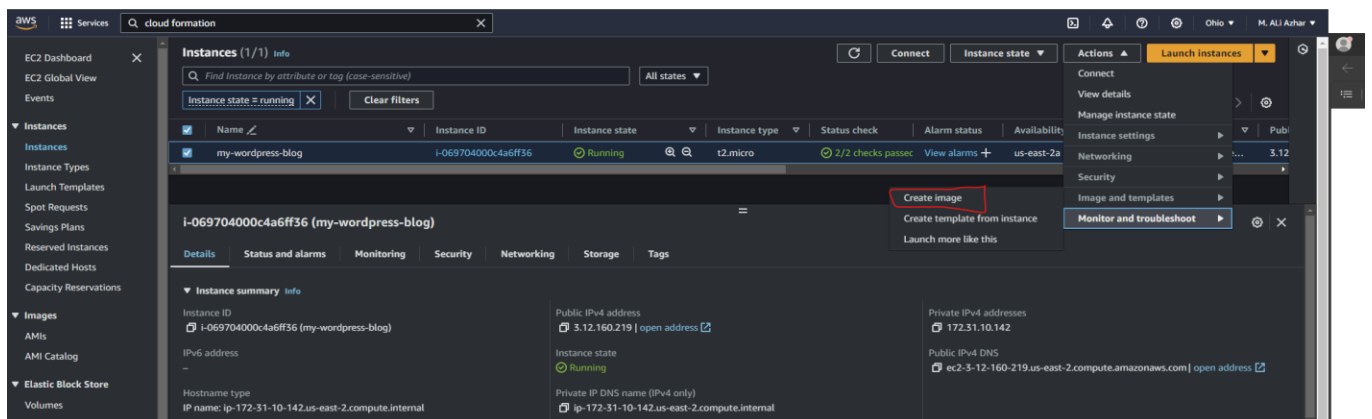


Fig 7: Creation of AMI from my new instance




Fig 8: Newly created Wordpress AMI status available

Configuration of Auto Scaling by using cloud formation:

First I create a cloud formation template with the below mentioned details in the yaml format.

1. Define the Resources:

- EC2 Instance: Use a T2.micro instance with our customized wordpress-blog AMI.
- Security Group: Configure the security group to allow HTTP (port 80) and SSH (port 22) access.
- Auto Scaling Group: Configure the Auto Scaling Group for the development environment.
- Launch Configuration: Create a launch configuration that uses my custom AMI.

2. CloudFormation Template Structure:

- **Parameters:** Define parameters for instance type, key pair name, and AMI ID.
- **Resources:** Define the EC2 instance, security group, and Auto Scaling Group.
- **Outputs:** Provide the public IP address of the WordPress instance.

3. Auto Scaling Group:

- I define an Auto Scaling Group that launches a new instance during business hours (9AM-6PM).
- Use a scheduled action in Auto Scaling to start and stop the development instance and set the desired capacity 2, Minimum capacity 1 and Maximum capacity 2.

After successfully completed and tested, I terminated all the instances and delete all resources to avoid any charges from AWS.

```

14
15 Resources:
16   EC2Instance:
17     Type: AWS::EC2::Instance
18     Properties:
19       KeyName: !Ref KeyName
20       InstanceType: "t2.micro"
21       ImageId: ami-090e8849fd96f3f
22       IamInstanceProfile: !Ref SessionManagerInstanceProfile
23       SecurityGroups:
24         - !Ref InstanceSecurityGroup
25
26   InstanceSecurityGroup:
27     Type: 'AWS::EC2::SecurityGroup'
28     Properties:
29       GroupDescription: Allow SSH access via port 22 and 80
30       SecurityGroupIngress:
31         - IpProtocol: tcp
32           FromPort: '22'
33           ToPort: '22'
34           CidrIp: !Ref SSHandWebLocation
35         - IpProtocol: tcp
36           FromPort: '80'
37           ToPort: '80'
38           CidrIp: !Ref SSHandWebLocation
39
40   SessionManagerRole:
41     Type: 'AWS::IAM::Role'
42     Properties:
43       AssumeRolePolicyDocument:
44         Version: 2012-10-17
45         Statement:
46           - Effect: Allow
47             Principal:
48               Service:
49                 - ec2.amazonaws.com
50             Action:
51               - 'sts:AssumeRole'
52       Path: /
53       ManagedPolicyArns:
54         - "arn:aws:iam::aws:policy/AmazonSSMManagedInstanceCore"

```

Fig 9: Cloud Formation Template for Auto-Scaling

```

MyLaunchConfig:
  Type: AWS::AutoScaling::LaunchConfiguration
  Properties:
    ImageId: ami-090e8849fd96f3f
    InstanceType: t2.micro
    SecurityGroups:
      - !Ref InstanceSecurityGroup
    KeyName: !Ref KeyName

MyAutoScalingGroup:
  Type: AWS::AutoScaling::AutoScalingGroup
  Properties:
    MinSize: 1
    MaxSize: 2
    DesiredCapacity: 2
    LaunchConfigurationName: !Ref MyLaunchConfig
    VPCZoneIdentifier:
      - subnet-0be53cae6cb78593d
      - subnet-0006b916f429b617a
      - subnet-05a7cb0959d70951c

ScaleOutSchedule:
  Type: AWS::AutoScaling::ScheduledAction
  Properties:
    AutoScalingGroupName: !Ref MyAutoScalingGroup
    DesiredCapacity: 1
    Recurrence: "0 9 * * *" # Every day at 9 AM

ScaleInSchedule:
  Type: AWS::AutoScaling::ScheduledAction
  Properties:
    AutoScalingGroupName: !Ref MyAutoScalingGroup
    DesiredCapacity: 0
    Recurrence: "0 18 * * *" # Every day at 6 PM

```

Fig 10: Cloud Formation Template to schedule AutoScaling

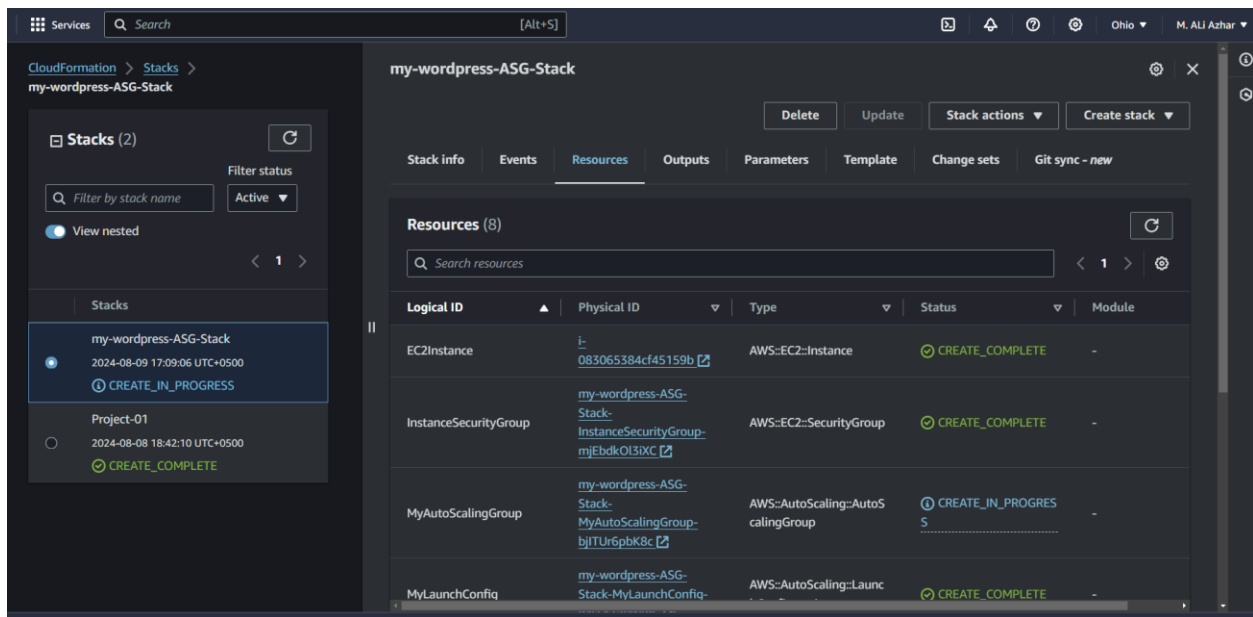


Fig 11: Autoscaling cloud formation stack creation

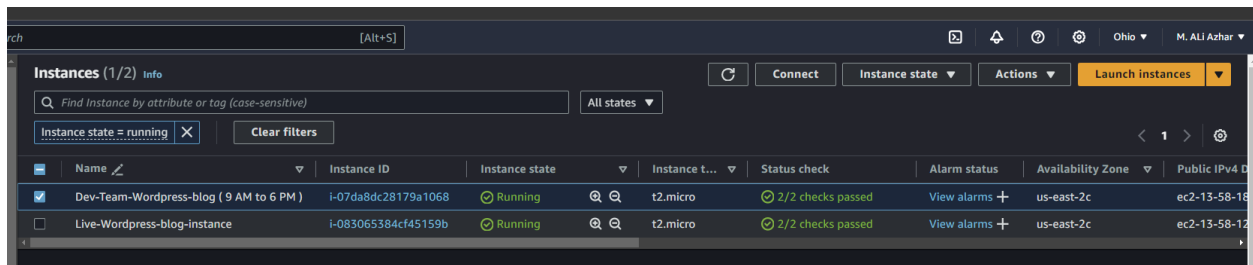


Fig 12: Dev instance created successfully at 9AM

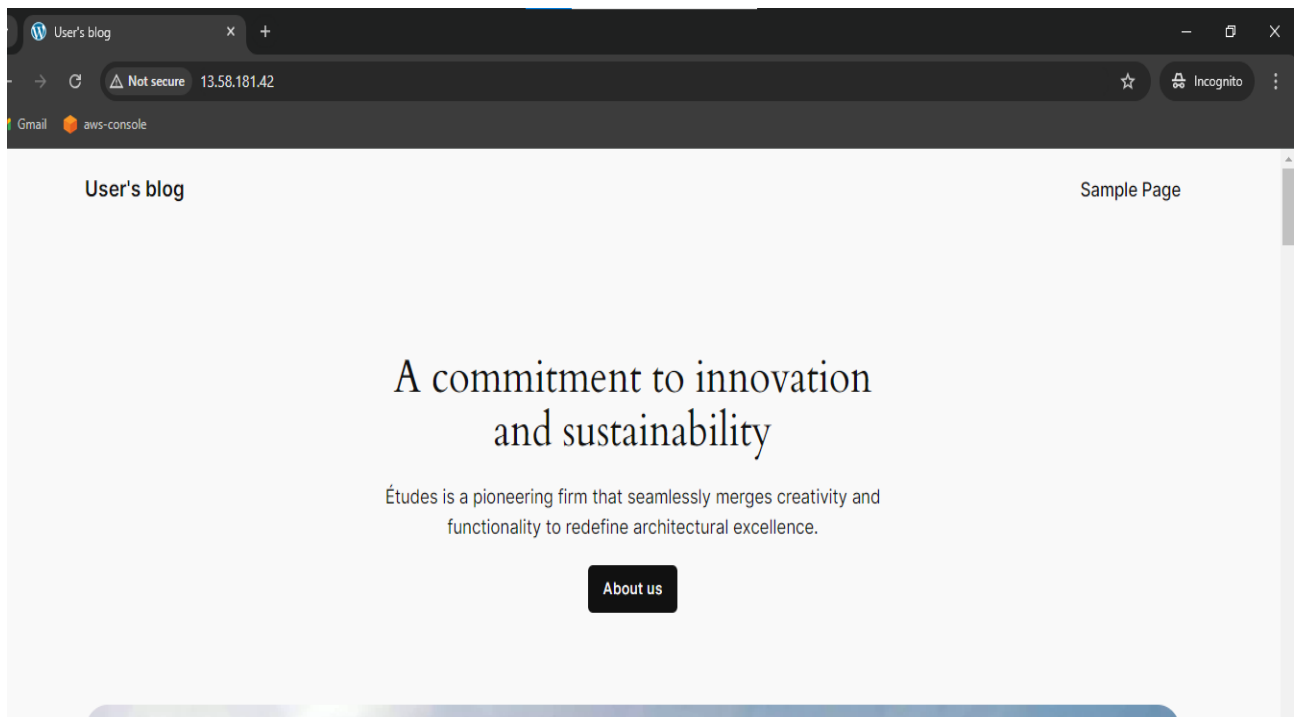


Fig 13: Dev instance sample page

The screenshot shows the AWS Management Console 'Instances' page. It displays a table of EC2 instances. The first instance, 'Dev-Team-Wordpress-blog', is in a 'Terminated' state. The second instance, 'Live-Wordpress-blog-instance', is in a 'Running' state. The table includes columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, and Public IPv4 address.

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 address
<input type="checkbox"/>	Dev-Team-Wordpress-blog (9 AM to 6 PM)	i-07da8dc28179a1068	Terminated	t2.micro	-	View alarms +	us-east-2c	-
<input type="checkbox"/>	Live-Wordpress-blog-instance	i-083065384cf45159b	Running	t2.micro	2/2 checks passed	View alarms +	us-east-2c	ec2-13-58-12...

Fig 14: Dev instance terminated at 6PM